

running example $H(x) = \sin(2x)(Z_0Z_1 + X_0 + X_1)$, $x = 0.7$, $T = 5.0 \mu\text{s}$, 2 qubits — ONE branch of each strategy, as emitted on the six physical channels

Nyquist mode $n=0$, shift $s=0.2067$ ($K=1.210$); machine residual $9.0\text{e-}05$ (source) / $8.8\text{e-}05$ (shifted). A gradient needs many branches either way — PSR spends them on a τ quadrature, NSR on drawn shifts; how many is Fig. 10's question, not this figure's.

